


Goal : Build a multi-document e-commerce Customer Support System.



Document Ingestion.

→ Loading the raw documents into the RAG pipeline.
↳ txt/pdf|.....

→ pypdf : Python library
↓
PdfReader

Process & Organize the document data.

↳ Raw documents are generally messy & contains some redundant data.

⇒ Parsing & Cleaning

Chunking

→ Searching through large documents is expensive.

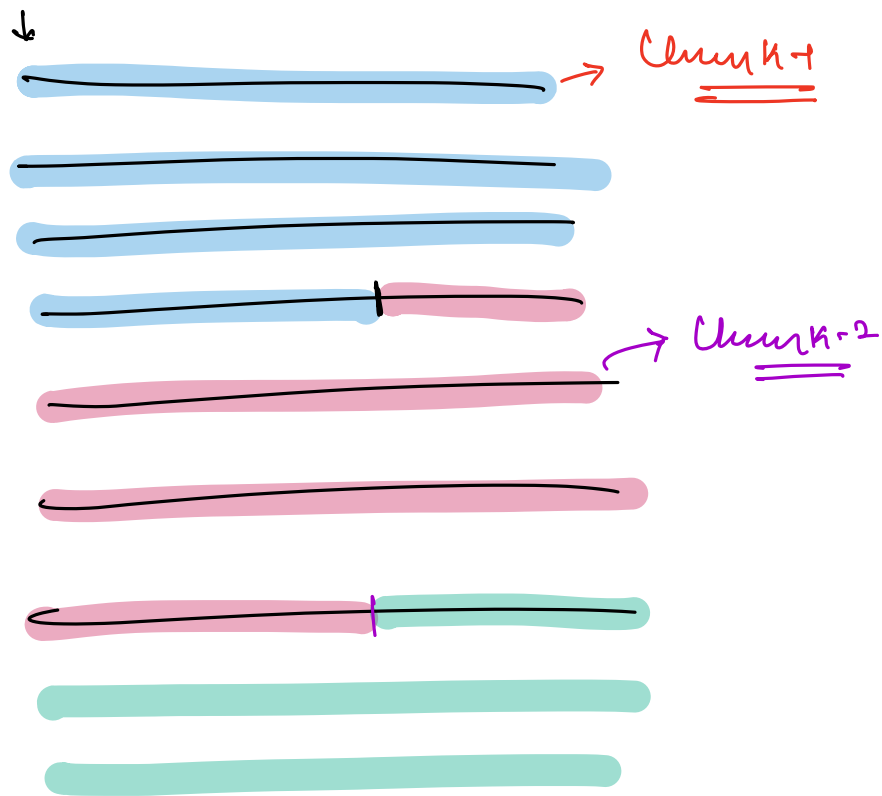
→ Sending large documents to the LLM is expensive.

↳ So, it's better to split documents into smaller chunks.

⇒ Chunking helps the retriever to find the relevant info more efficiently.

Fixed Size Chunking

ex: Chunk Size = 100 words.



⇒ Some context will get lost.

→ Semantic Chunking

→ Hierarchical Chunking

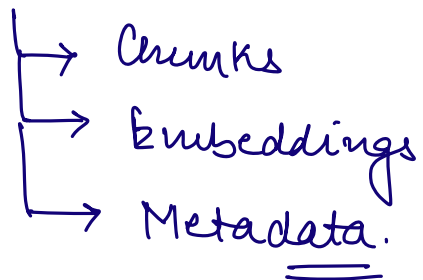
Chunk Overlap

Embeddings.

→ After chunking, we convert each chunk into embedding.

Embedding Model : text-embedding-3-small → 1536
text-embedding-3-large → 3072

Vector DB.



Retriever

