

End-to-end Implementation workflow:

I) Data.

↳ Documents | Files | Text | Pdf - - - -

II) Embeddings

↳ Convert each data item into a vector embedding using some model.

III) Storage.

↳ Store embeddings into Vector DB.


- Original text
- Vector embedding
- Unique Id
- metadata.

IV) Query


↳ User asks a query & we convert this query into embedding.

V) Similarity Search.

↳ Search in vector DB and find top-k nearest neighbors wrt the query embedding.



Development Environment Setup.



Vector DB → Chroma DB.

↓

Good choice for local implementation.

```
graph TD; Chroma[Chroma DB] --> Choice[Good choice for local implementation];
```

⇒

1. Sample documents.
2. Generate embeddings.
3. Store these embeddings into Chroma DB.
4. Convert the query into embedding
5. Perform Similarity Search.