

Pre-trained model
[It is trained on huge amount
of internet data]

⇒ Hallucination.

↳ Not desirable
↳ Try to reduce hallucination as much as
possible.

⇒ LLMs are very powerful but they have limited
knowledge.

limitations of LLM:

→ Static knowledge.

→ Doesn't have access to private policy documents.

→ Hallucination.

Grounded AI Systems.

Model's response should be based on external, verifiable, reliable information & not only on internal memory.

RAG.

↳ Retrieval Augmented Generation.

It is an approach where the system:

→ takes the user query

→ retrieves the information from external knowledge sources.

→ Provide this information to the LLM as context.

→ LLM generates the answer using the retrieved context.

RAG combines retrieval of relevant information from external sources & generation of responses in natural language.

So, instead of answering only from memory, the system answers using external sources as well.

RAG helps solve the main limitations of plain LLM.

"Grounding" in RAG.


→ Reduces Hallucinations.

→ More reliable, accurate & trustworthy answers.

→ Makes it easier to trace the source of the information.

RAG flow.

⇒ We have already stored data in the form of embeddings in Vector DB.



#	LLM-Only	RAG Based
Knowledge Source	Only pre-trained memory	pre-trained memory + external <u>source</u>
Access to latest info	Usually <u>No</u>	<u>Yes</u>
Access to Private info	No	<u>Yes</u>
Hallucination	High	Very <u>low</u> .

RAG + Vector DB.


⇒ RAG based applications use Vector DB to store & retrieve information.

Exact Match

Similarity Search.

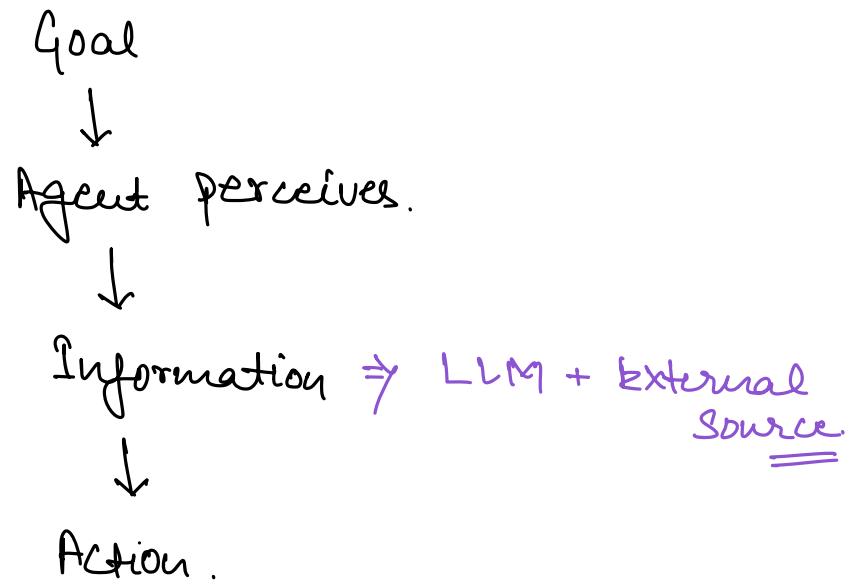
⇒ Modern AI applications use combination of Exact match & Similarity search for better responses.

#



RAG in Agentic AI.

→ Agentic Flow



the informal