


⇒ RAQ application using LangChain.

↳ Retrieval Augmented Generation.

LLMs are very powerful but mainly there's one limitation:

↳ Static pre-trained knowledge.

RAQ.



User Query



Retriever

[Searches relevant documents]



Relevant Chunks are
given to the LLM to
generate the final ans.

1) Document loader.

→ Reads the external data & converts it into LangChain Document Objects.

```
Document(  
    page-content = " — " ,  
    metadata = " — "  
)
```

2) Text Splitter

It is difficult to retrieve the data from bigger documents, that's why we split the documents into smaller chunks.

RecursiveCharacterTextSplitter

→ chunk-size

→ chunk-overlap

3) Embedding

word
↓

[- - - - -]

OpenAI Embedding model

- small ⇒ 1536
- large ⇒ 3072

Close Vectors $\Rightarrow$ Similar Meanings.

4) Vector DB

→ Store embeddings & allows Similarity Search.
→ Chroma DB.

5) Retriever

→ Takes the user query & searches relevant documents.

LCB L

retriever | format doc | prompt | LLM | parser